

Time Series Analysis

Stationarity, differencing and linear modeling

Concepts

- ▶ Elements of exploratory time series analysis
 - ▶ Stationarity and differencing
- ▶ Models of time series
 - ▶ Linear models and stochastic processes




What is a stationary time series?


- ▶ “...a time series is said to be stationary if there is no systematic change in mean (no trend), if there is no systematic change in variance and if strictly periodic variations have been removed.” (Chatfield: 13)
- ▶ Correlogram has very few significant spikes at very small lags and cuts off drastically/dies down quickly (at 2 or 3 lags)



Example: Stationary series



The correlogram has only a couple of significant spikes at lags 1 and 2
So the correlogram confirms that the series is stationary



Why do we care about stationarity?

- ▶ Most time series models only work if data are stationary
 - ▶ Called “stationary time series models”
- ▶ So, first step in an analysis is to check for evidence of a trend or seasonality
- ▶ Depending on the type of trend or seasonality, we can remove nonstationarity in different ways
 - ▶ Simple differencing subtracts trend/seasonal component from original series
 - ▶ Regression can model trend/seasonal components
 - ▶ Use the time series of residuals in model



Stochastic vs. deterministic


- ▶ Used to describe the “quality” of nonstationarity of a series
 - ▶ Help determine what type of TS approach to take
- ▶ Stochastic = inexplicable changes in direction
 - ▶ Often found in economic processes, sometimes climate
 - ▶ “Random walk” process
 - ▶ Use differencing and autoregressive models
- ▶ Deterministic = plausible physical explanation for a trend or seasonal cycle
 - ▶ Increase in population, orbit of the earth
 - ▶ Use regression models



Example: stochastic TS



Example: deterministic TS



To make a series stationary

1. Check if there is variance that changes with time
 - ▶ YES → make variance constant with log or square root transformation
2. Is the trend stochastic or deterministic?
 - ▶ If stochastic → use differencing
 - ▶ If deterministic → use regression
3. Remove the trend in mean with:
 - ▶ 1st/2nd order differencing
 - ▶ Smoothing and differencing (seasonality)
4. If there is seasonality in the data:
 - ▶ Moving average and differencing
 - ▶ Smoothing



Differencing

Making non-stationary time series stationary
Stochastic seasonality and trends

- ▶ If the series has a stable long-run trend and tends to revert to the trend line following a disturbance
 - ▶ **trend-stationary**
- ▶ Sometimes even de-trending is not sufficient to make the series stationary, in which case it may be necessary to transform it into a series of period-to-period and/or season-to-season *differences*.
- ▶ If the mean, variance, and autocorrelations of the original series are not constant in time, even after detrending, perhaps the statistics of the *changes* in the series between periods or between seasons *will* be constant. Such a series is said to be **difference-stationary**.

Differencing

- ▶ Transformation of the series to a new time series where the values are the differences between consecutive values
 - ▶ Procedure may be applied consecutively more than once, giving rise to the "first differences", "second differences", etc.

- ▶ The first order differences are computed as :

$$d^1_t = x_t - x_{t-1}$$


The second order differences can be computed as :

$$d^2_t = d^1_t - d^1_{t-1}$$

Seasonal differences can be computed as :

$$d^{12}_t = d_t - d_{t-12}$$

US Unemployment



Checking for stationarity

```
> install.packages("tseries")
```

```
> library(tseries)
```

```
> kpss.test(US.ts, null = "Trend")
```

KPSS Test for Trend Stationarity

data: US.ts

KPSS Trend = 0.5118, Truncation lag parameter = 2, p-value = 0.01

Can reject H_0 - means the series is non-stationary

```
> adf.test(US.ts, alternative = "stationary") # can also add k=n
```

Augmented Dickey-Fuller Test


data: US.ts

Dickey-Fuller = -1.2823, Lag order = 5, p-value = 0.8749

Cannot reject H_0 - means the series is non-stationary


Example

```
> acf(US.ts)
> pacf(US.ts)
```




Example

```
> dUSts<-diff(US.ts)
> plot(dUSts)
```



Example



Checking for stationarity

```
> kpss.test(dUSts, null = "Trend")
```

KPSS Test for Trend Stationarity

data: dUSts

KPSS Trend = 0.3197, Truncation lag parameter = 2, p-value = 0.01

```
> adf.test(dUSts, alternative = "stationary")
```


Augmented Dickey-Fuller Test

data: dUSts

Dickey-Fuller = -3.059, Lag order = 5, p-value = 0.1363

Example

```
> ddUSts<-diff(dUSts)
```



Example

```
> kpss.test(ddUSts, null = "Trend")
```

KPSS Test for Trend Stationarity

data: ddUSts

KPSS Trend = 0.0139, Truncation lag parameter = 2, p-value = 0.1

```
> adf.test(ddUSts, alternative = "stationary")
```

Augmented Dickey-Fuller Test

data: ddUSts


Dickey-Fuller = -8.2463, Lag order = 5, p-value = 0.01

Seasonality and smoothing

- ▶ Smoothing is often used to remove an underlying signal or trend (such as a seasonal cycle)
- ▶ Common method is the centered moving average
 - ▶ Average a specified number of time series values around each value in the time series
 - ▶ Length of moving average is chosen to average out seasonal effects
 - ▶ Seasonal = 12 point moving average
 - ▶ Quarterly = 4 point moving average




Example: river flow rates




Example

```
> mbc.d<-decompose(mbc)
> dmbc<-(mbc.d$trend+mbc.d$rand)
> plot(dmbc)
```



Example

```
> acf(dmbc[7:582]) # moving averages leave out first/last 6 obs
```



What now? Modeling a stationary series

- ▶ Once you have a stationary series, you can model it using an autoregressive (AR) model

$$x_t = \beta_1 x_{t-1} + \dots + \beta_p x_{t-p} + r_t$$

where :

x time series observation

p process or the order (lag)

r random element (error)

- ▶ We simply regress x on lagged x
- ▶ If the model successfully captures the dependence structure in the data the residuals should look random
 - ▶ Check the residuals from the AR model for any “left-over” dependence



Autoregressive model

```
> mbc.ar<-ar(dmbc[7:582], method="mle")
> mean(dmbc[7:582])
[1] 1.530268
```

```
> mbc.ar
```

Call:

```
ar(x = dmbc[7:582], method = "mle")
```

Coefficients:


```
1
0.3896
```

$$z_t = 1.55 + 0.38(z_{t-1} - 1.55)$$

```
Order selected 1 sigma^2 estimated as 0.6637
```

```
> plot(mbc.ar$res); acf(mbc.ar$res[7:582]); pcf(mbc.ar$res[7:582])
```





Regression


Making non-stationary time series stationary
Working with deterministic time series

Regression

- ▶ **Deterministic trends can be modeled using regression techniques**
 - ▶ Both linear and nonlinear
- ▶ **Seasonal trends can be modeled using regression techniques**
 - ▶ Indicator variables
 - ▶ Harmonic variables
- ▶ **Residuals from models are the “random error” component**
 - ▶ Use these for time series models

▶

Example: Mona Loa CO₂ Concentrations



▶

Fitting linear models to TS data

- ▶ Regress observation against time
 - $X_t = t$
 - ▶ Typically use generalized least squares (GLS) because errors are correlated
 - ▶ Type of maximum likelihood method
- ▶ Examine coefficients and SE
- ▶ Examine correlogram of residuals
- ▶ Fit appropriate TS model to residuals



Example

```
> mllm<-lm(mlco2~mlco2.t)
> acf(resid(mllm))$acf[2]
[1] 0.9261262

> mlnlm<-gls(mlco2~mlco2.t, cor=corAR1(0.93))
> mlnlm
Generalized least squares fit by REML
Model: mlco2 ~ mlco2.t
Log-restricted-likelihood: -993.7159
```

Coefficients:

(Intercept)	mlco2.t
-2488.298538	1.428678

Degrees of freedom: 623 total; 621 residual

Residual standard error: 3.551621

```
> confint(mlnlm)
```

	2.5 %	97.5 %
(Intercept)	-2686.601893	-2289.995183
mlco2.t	1.328734	1.528622

CI's for slope do not encompass 0

1) Estimates are significant

2) Trend is significant



- ▶ To take account of seasonal effects, add predictor variables for season
- ▶ Basically, a set of dummy variables that indicate the season, quarter, month, etc.

$$X_t = \alpha_1 T + S_t + z_t$$

Creating seasonal indicators

```
> Seas<-cycle(mlco2)
> Seas
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1958			3	4	5	6	7	8	9	10	11	12
1959	1	2	3	4	5	6	7	8	9	10	11	12
1960	1	2	3	4	5	6	7	8	9	10	11	12
1961	1	2	3	4	5	6	7	8	9	10	11	12

```
> Time<-time(mlco2)
> Time
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	...
1958			1958.167	1958.250	1958.333	1958.417	1958.500	1958.583	...
1959	1959.000	1959.083	1959.167	1959.250	1959.333	1959.417	1959.500	1959.583	...
1960	1960.000	1960.083	1960.167	1960.250	1960.333	1960.417	1960.500	1960.583	...
1961	1961.000	1961.083	1961.167	1961.250	1961.333	1961.417	1961.500	1961.583	...
1962	1962.000	1962.083	1962.167	1962.250	1962.333	1962.417	1962.500	1962.583	...



Results


```
> ml.gls<-glsl(mlco2~Time+factor(Seas), cor=corAR1(0.985))
```

Variable	Coefficient
(Intercept)	-2492.85
Time	1.43
factor(Seas)2	0.63
factor(Seas)3	1.39
factor(Seas)4	2.51
factor(Seas)5	2.92
factor(Seas)6	2.27
factor(Seas)7	0.67
factor(Seas)8	-1.45
factor(Seas)9	-3.13
factor(Seas)10	-3.26
factor(Seas)11	-2.11
factor(Seas)12	-0.94




Results

```
> acf(ml.gls$residuals)
> pacf(ml.gls$residuals)
```



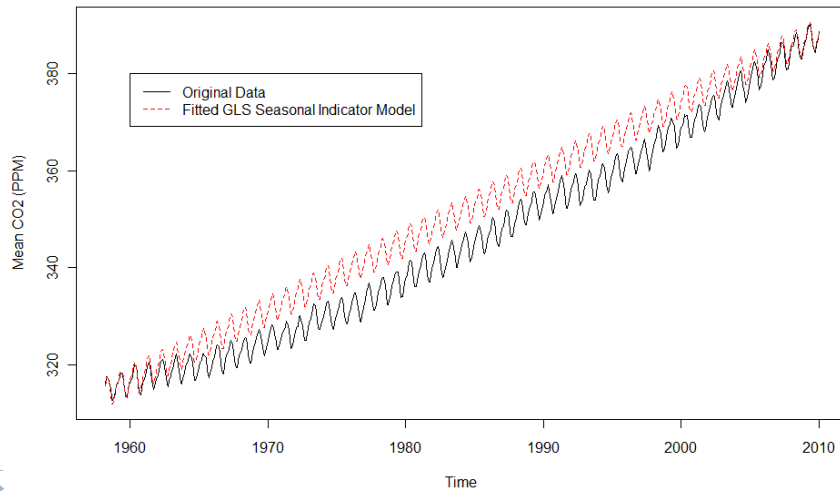
Results

```
> ts.plot(ml.gls$residuals, ylab="GLS Residuals")
```



Results

```
> ts.plot(cbind(mlco2, ml.gls$fitted), lty=1:2, col=c(1,2))
> legend(1960, 380, c("Original", "Fitted"), col=c(1,2), lty=c(1, 2))
```



Results

```
> ml<-lm(mlco2~mlco2.t)
> AIC(ml)
[1] 3248.806

> ml.lm<-lm(mlco2~Time+factor(Seas))
> AIC(ml.lm)
[1] 2958.362

> ml.gls<-glS(mlco2~Time+factor(Seas), cor=corAR1(0.985))
> AIC(ml.gls)
[1] 391.0092
```


Seasonal harmonic model

- ▶ Rational is that seasonal effects vary smoothly over the seasons
- ▶ Instead of using an indicator of season (1=January, 0=Not), can use a smooth function
 - ▶ Sine and/or cosine wave that oscillates at a certain number of cycles per season


$$X_t = T_t + \sum_{i=1}^{s/2} \{s_i \sin(2\pi i t / s) + c_i \cos(2\pi i t / s)\} + z_t$$



One sine



Two sines, Two cosines



Example

- First, we need to simulate a set of empty sine and cosine waves at different frequencies (1-6)

```
> SIN<-COS<-matrix(nr=length(mlco2), nc=6)
> for (i in 1:6) {
  COS[,i] <-cos(2*pi*i*time(mlco2))
  SIN[,i] <-sin(2*pi*i*time(mlco2)) }
```


- We don't know how many harmonics we need to include
- Fortunately, harmonic coefficients are known to be independent
 - Can add them all in, check for statistical significance and only keep the ones that are



Example

```
> TIME <- (time(mlco2) - mean(time(mlco2))) / sd(time(mlco2))
> mean(time(mlco2))
> sd(time(mlco2))

> ml.shgls<-gls(mlco2~ TIME + I(TIME^2) +
  COS[,1]+SIN[,1]+COS[,2]+SIN[,2]+
  COS[,3]+SIN[,3]+COS[,4]+SIN[,4]+
  COS[,5]+SIN[,5]+COS[,6]+SIN[,6], corr=corAR1(0.908))
```



(Intercept)	TIME	I (TIME^2)	COS[, 1]
1600.58752636	155.24459467	18.22074618	-9.27557755
SIN[, 1]	COS[, 2]	SIN[, 2]	COS[, 3]
80.77815539	19.75028133	-37.63167836	-2.19576028
SIN[, 3]	COS[, 4]	SIN[, 4]	COS[, 5]
-6.27532860	2.24416895	5.78236098	0.97726215
SIN[, 5]	COS[, 6]	SIN[, 6]	
2.60128993	0.02081744	0.64617844	



Results


```
> ml.shgls2<-gls(mlco2~ TIME + I (TIME^2) +
  COS[,1]+SIN[,1]+COS[,2]+SIN[,2]+
  COS[,3]+SIN[,3]+COS[,4]+SIN[,4]+
  SIN[,5], corr=corAR1(0.908))
> coef(ml.shgls2)/sqrt(diag(vcov(ml.shgls2)))
(Intercept)      TIME      I (TIME^2)      COS[, 1]      SIN[, 1]
1599.340883  155.135181  18.209209    -9.296485    80.892083
      COS[, 2]      SIN[, 2]      COS[, 3]      SIN[, 3]      COS[, 4]
19.878676   -37.752959   -2.203375    -6.278049     2.251647
      SIN[, 4]      SIN[, 5]
      5.792330     2.595673

> AIC(ml.shgls2)       $x_t = 1599 + 155.1t + 18.2t^2 - 9.3\cos(2\pi/12) + 80.9\sin(2\pi/12) +$ 
[1] 367.8002            $19.9\cos(4\pi/12) - 37.8\sin(4\pi/12) - 2.2\cos(6\pi/12) + \dots +$ 
                        $2.6\sin(10\pi/12) + z_t$ 
```


Plot residuals and acf/pacf for ml.shgls2 residuals




Results



Results




Results



Now the autoregressive model

```
> ml.ar<-ar(resid(ml.shgls2), order=2)
> plot(ml.ar$resid, ylab="Residuals", type="p"); abline(h=0)
> acf(ml.ar$resid[-(1:2)], main="")
```



What's the equation?

$$x_t = 1599 + \frac{155.1(t-1984)}{15} + \frac{18.2(t-1984)^2}{15} - 9.3 \cos(2\pi t / 12) + 80.9 \sin(2\pi t / 12) + 19.9 \cos(4\pi t / 12) - 37.8 \sin(4\pi t / 12) - 2.2 \cos(6\pi t / 12) - 6.3 \sin(6\pi t / 12) + 2.3 \cos(8\pi t / 12) + 5.8 \sin(8\pi t / 12) + 2.6 \sin(10\pi t / 12) + z_t$$

Annotations: Red arrows point from the text 'mean of t' to the linear term coefficient and from 'SD of t' to the quadratic term coefficient. The residual term z_t is enclosed in a red box.

where the residual series follows an AR(2) process:

$$z_t = .72e_{t-1} + .21e_{t-2} + e_t$$



Making predictions


- ▶ The generic way of making predictions is `predict()`
 - ▶ But, must create a new data frame with values used to predict properly labeled

```
> new.T<-time(ts(start=2010, end=c(2020, 12), fr=12))
> TIME<-(new.T-mean(time(mlco2)))/sd(time(mlco2))
> SIN<-COS<-matrix(nr=length(new.T), nc=6)
> for (i in 1:6) {
  COS[,i] <-cos(2*pi*i*time(new.T))
  SIN[,i] <-sin(2*pi*i*time(new.T)) }
> SIN<-SIN[,-6]
> COS<-COS[,-(5:6)]
> new.dat<-data.frame(TIME=as.vector(TIME), SIN=SIN, COS=COS)
> ml.pred.ts<-ts(predict(ml.shgls2, new.dat), st=2010, fr=12)
```



Results

```
> plot(mlco2, ml.pred.ts, lty=1:2, col=c(1,2))
```



What have we learned?

- ▶ The importance of stationarity
- ▶ The difference between stochastic and deterministic trends in the data
- ▶ Methods for making stochastic TS stationary
 - ▶ Differencing
 - ▶ Smoothing
- ▶ Regression techniques for modeling deterministic TS
 - ▶ Removes both trend and seasonality



What's next?

- ▶ Examining and identifying stochastic processes that give rise to an observed time series
 - ▶ Autoregressive processes (AR)
 - ▶ How strongly the past influences the present
 - ▶ Moving average processes (MA)
 - ▶ Model present using past errors of prediction
 - ▶ Autoregressive moving average processes (ARMA)

